

RAGAS Metrics Guide

Complete guide to evaluating RAG applications

Overview

RAGAS (Retrieval Augmented Generation Assessment) provides metrics to evaluate both the retriever and the LLM in your RAG pipeline.

- **Retriever Metrics:** Context Precision, Context Recall, Context Entities Recall
 - **LLM Metrics:** Answer Relevancy, Faithfulness, Noise Sensitivity
-

1. Answer Relevancy

What is it?

Measures if the chatbot's answer actually addresses the user's question.

Score: 0-1 (higher = better)

How it Works

1. LLM generates 3 questions from the answer
2. Compares with original question using cosine similarity
3. Averages scores = final score

What You Provide

- `user_input`: The question asked
- `response`: Chatbot's answer

Code Example

```
from ragas.metrics import AnswerRelevancy

scorer = AnswerRelevancy(llm=llm)
result = await scorer.ascore(
    user_input="What is your refund policy?",
    response="Refunds available within 30 days."
)
```

Example: HIGH Score

Q: How to cancel flight? | **A:** Cancel via app under My Trips.

Generated Qs match original → **Score: 0.92**

Example: LOW Score

Q: How to cancel flight? | **A:** Flynas offers meals on flights.

Generated Qs don't match → **Score: 0.15**

2. Faithfulness

What is it?

Measures if the answer is factually correct based on retrieved context.

Score: 0-1 (higher = more faithful)

How it Works

1. Breaks answer into individual statements
2. Checks if each statement is supported by context
3. Score = (supported statements) / (total statements)

What You Provide

- `response`: Chatbot's answer
- `retrieved_contexts`: Documents the RAG retrieved

Code Example

```
from ragas.metrics import Faithfulness

scorer = Faithfulness(llm=llm)
result = await scorer.ascore(
    response=test_case["response"],
    retrieved_contexts=test_case["retrieved_contexts"]
)
```

Example: HIGH Score

Context: Free cancel within 24hrs | **Answer:** Cancel free within 24hrs.

All statements supported → **Score: 1.0**

Example: LOW Score (Hallucination)

Context: Free cancel within 24hrs | **Answer:** Cancel free anytime.

Statement NOT in context → **Score: 0.0**

3. Context Precision

What is it?

Measures if the retriever ranks relevant documents at the top.

Score: 0-1 (higher = better ranking)

How it Works

1. Checks each retrieved document for relevance
2. Calculates precision at each position
3. Relevant docs at top = HIGH, at bottom = LOW

What You Provide (Without Reference Version)

- `user_input`: The question
- `response`: The chatbot's answer
- `retrieved_contexts`: Documents in order

Code Example (Without Reference - Recommended)

```
import asyncio
from ragas.metrics import LLMContextPrecisionWithoutReference
from ragas import SingleTurnSample
from ragas.llms import LangchainLLMWrapper
from langchain_openai import ChatOpenAI

async def test():
    llm = LangchainLLMWrapper(ChatOpenAI(model="gpt-4o"))
    scorer = LLMContextPrecisionWithoutReference(llm=llm)

    sample = SingleTurnSample(
        user_input="What is baggage allowance?",
        response="Economy allows 23kg checked baggage.",
        retrieved_contexts=[
            "Economy class allows 23kg checked baggage.", # relevant
            "Flynas serves meals on flights.", # irrelevant
        ]
    )

    score = await scorer.single_turn_ascore(sample)
    print(f"Score: {score}")

asyncio.run(test())
```

Example: HIGH Score

Retrieved: [Cancellation policy (relevant), Baggage, Meals]

Relevant doc at position 1 → **Score:** ~1.0

Example: LOW Score

Retrieved: [Meals, Baggage, Cancellation policy (relevant)]

Relevant doc at position 3 → **Score:** ~0.3

4. Context Recall

What is it?

Measures if the retriever fetched ALL relevant documents needed.

Score: 0-1 (higher = more complete)

How it Works

1. Breaks reference into individual claims
2. Checks if each claim is in retrieved contexts
3. Score = (claims found) / (total claims)

What You Provide

- `retrieved_contexts`: Documents retrieved
- `reference`: Ground truth answer (required!)

Code Example

```
from ragas.metrics import LLMContextRecall

scorer = LLMContextRecall(llm=llm)
result = await scorer.ascore(
    retrieved_contexts=test_case["retrieved_contexts"],
    reference=test_case["reference"]
)
```

Example: HIGH Score

Reference has 2 claims, both found in retrieved docs.

2/2 claims found → **Score: 1.0**

Example: LOW Score

Reference has 2 claims, only 1 found in retrieved docs.

1/2 claims found → **Score: 0.5**

Precision vs Recall

- **Precision:** Are relevant docs at TOP?
 - **Recall:** Did we get ALL relevant docs?
-

5. Context Entities Recall

What is it?

Measures if retriever fetched docs with all important entities (names, dates, places).

Score: 0-1 (higher = more entities)

How it Works

1. Extracts entities from reference (names, dates, places)
2. Checks how many appear in retrieved contexts
3. Score = (entities found) / (total entities)

What You Provide

- `retrieved_contexts`: Documents retrieved
- `reference`: Ground truth with entities

Code Example

```
from ragas.metrics import ContextEntityRecall

scorer = ContextEntityRecall(l1m=l1m)
result = await scorer.ascore(
    retrieved_contexts=test_case["retrieved_contexts"],
    reference=test_case["reference"]
)
```

Example: HIGH Score

Reference: [Flynas, Flight 123, Riyadh, 10:00 AM] = 4 entities

All 4 found in docs → **Score: 1.0**

Example: LOW Score

Reference: [Flynas, Flight 123, Riyadh, 10:00 AM] = 4 entities

Only Flynas found → **Score: 0.25**

6. Noise Sensitivity

What is it?

Measures how often LLM makes errors when given noisy/irrelevant docs.

Score: 0-1 (LOWER = better, opposite of others!)

How it Works

1. Checks each claim in LLM response
2. Verifies if correct against ground truth
3. Score = (incorrect claims) / (total claims)

What You Provide

- `user_input`: Question
- `response`: LLM answer
- `reference`: Ground truth
- `retrieved_contexts`: Docs including noise

Code Example

```
from ragas.metrics import NoiseSensitivity

scorer = NoiseSensitivity(llm=llm)
result = await scorer.ascore(
    user_input=test_case["question"],
    response=test_case["response"],
    reference=test_case["reference"],
    retrieved_contexts=test_case["retrieved_contexts"]
)
```

Example: LOW Score (GOOD)

Retrieved: [Baggage policy, Meal menu (noise)] | **Answer:** Only about baggage.

LLM ignored noise → **Score: 0.0 (GOOD)**

Example: HIGH Score (BAD)

Retrieved: [Baggage policy, Meal menu (noise)] | **Answer:** Baggage + meals.

LLM confused by noise → **Score: 0.5 (BAD)**

Quick Reference Summary

Metric Comparison

Metric	What it Tests	Tests
Answer Relevancy	Does answer address the question?	LLM
Faithfulness	Is answer factually correct from context?	LLM
Context Precision	Are relevant docs ranked at top?	Retriever
Context Recall	Did we fetch ALL relevant docs?	Retriever
Context Entities Recall	Did we fetch all entities?	Retriever
Noise Sensitivity	Does LLM resist bad docs?	LLM (Lower=Better)

What Each Metric Needs

Metric	Required Inputs
Answer Relevancy	<code>user_input, response</code>
Faithfulness	<code>response, retrieved_contexts</code>
Context Precision (without ref)	<code>user_input, response, retrieved_contexts</code>
Context Recall	<code>retrieved_contexts, reference</code>
Context Entities Recall	<code>retrieved_contexts, reference</code>
Noise Sensitivity	<code>user_input, response, reference, retrieved_contexts</code>

Key Insight

To use most metrics, you need to capture `retrieved_contexts` from your RAG pipeline. Modify your chatbot code to log the documents fetched during retrieval.

References

- RAGAS Docs - Answer Relevancy
- RAGAS Docs - Faithfulness
- RAGAS Docs - Context Precision
- RAGAS Docs - Context Recall